

Ruby on Rails Tutorial

Learn Rails by Example

Michael Hartl

Foreword

My former company (CD Baby) was one of the first to loudly switch to Ruby on Rails, and then even more loudly switch back to PHP (Google me to read about the drama). This book by Michael Hartl came so highly recommended that I had to try it, and *Ruby on Rails Tutorial* is what I used to switch back to Rails again.

Though I’ve worked my way through many Rails books, this is the one that finally made me “get” it. Everything is done very much “the Rails way”—a way that felt very unnatural to me before, but now after doing this book finally feels natural. This is also the only Rails book that does test-driven development the entire time, an approach highly recommended by the experts but which has never been so clearly demonstrated before. Finally, by including Git, GitHub, and Heroku in the demo examples, the author really gives you a feel for what it’s like to do a real-world project. The tutorial’s code examples are not in isolation.

The linear narrative is such a great format. Personally, I powered through *Rails Tutorial* in three long days, doing all the examples and challenges at the end of each chapter. Do it from start to finish, without jumping around, and you’ll get the ultimate benefit.

Enjoy!

Derek Sivers (sivers.org)

Formerly: Founder, [CD Baby](#)

Currently: Founder, [Thoughts Ltd.](#)

Acknowledgments

Ruby on Rails Tutorial owes a lot to my previous Rails book, *RailsSpace*, and hence to my coauthor [Aurelius Prochazka](#). I’d like to thank Aure both for the work he did on that book and for his support of this one. I’d also like to thank Debra Williams Cauley, my editor on both *RailsSpace* and *Rails Tutorial*; as long as she keeps taking me to baseball games, I’ll keep writing books for her.

I’d like to acknowledge a long list of Rubyists who have taught and inspired me over the years: David Heinemeier Hansson, Yehuda Katz, Carl Lerche, Jeremy Kemper, Xavier Noria, Ryan Bates, Geoffrey Grosenbach, Peter Cooper, Matt Aimonetti, Gregg Pollack, Wayne E. Seguin, Amy Hoy, Dave Chelimsky, Pat Maddox, Tom Preston-Werner, Chris Wanstrath, Chad Fowler, Josh Susser, Obie Fernandez, Ian McFarland, Steven Bristol, Giles Bowkett, Evan Dorn, Long Nguyen, James Lindenbaum, Adam Wiggins, Tikhon Bernstam, Ron Evans, Wyatt Greene, Miles Forrest, the good people at Pivotal Labs, the Heroku gang, the thoughtbot guys, and the GitHub crew. Finally, many, many readers—far too many to list—have contributed a huge number of bug reports and suggestions during the writing of this book, and I gratefully acknowledge their help in making it as good as it can be.

About the author

[Michael Hartl](#) is a programmer, educator, and entrepreneur. Michael was coauthor of *RailsSpace*, a best-selling Rails tutorial book published in 2007, and was cofounder and lead developer of [Insoshi](#), a [popular](#) social networking platform in Ruby on Rails. Previously, he taught theoretical and computational physics at the [California Institute of Technology](#) (Caltech), where he received the Lifetime Achievement Award for Excellence in Teaching. Michael is a graduate of [Harvard College](#), has a [Ph.D. in Physics](#) from [Caltech](#), and is an alumnus of the [Y Combinator](#) program.

Copyright and license

Ruby on Rails Tutorial: Learn Rails by Example. Copyright © 2010 by Michael Hartl. All source code in *Ruby on Rails Tutorial* is available under the [MIT License](#) and the [Beerware License](#).

Copyright (c) 2010 Michael Hartl

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

```
/*
 * -----
 * "THE BEERWARE LICENSE" (Revision 42):
 * Michael Hartl wrote this code. As long as you retain this notice, you can
 * do whatever you want with this stuff. If we meet someday, and you think
 * this stuff is worth it, you can buy me a beer in return.
 * -----
 */
```

Chapter 4 Rails-flavored Ruby

Grounded in examples from [Chapter 3](#), this chapter explores some elements of Ruby important for Rails. Ruby is a big language, but fortunately the subset needed to be productive as a Rails developer is relatively small. Moreover, this subset is *different* from the usual approaches to learning Ruby, which is why, if your goal is making dynamic web applications, I recommend learning Rails first, picking up bits of Ruby along the way. To be a Rails *expert*, you need to understand Ruby more deeply, and this

book gives you a good foundation for developing that expertise. As noted in [Section 1.1.1](#), after finishing *Rails Tutorial* I suggest reading a pure Ruby book such as [Beginning Ruby](#), [The Well-Grounded Rubyist](#), or [The Ruby Way](#).

This chapter covers a lot of material, and it's OK not to get it all on the first pass. I'll refer back to it frequently in future chapters.

4.1 Motivation

As we saw in the last chapter, it's possible to develop the skeleton of a Rails application, and even start testing it, with essentially no knowledge of the underlying Ruby language. We did this by relying on the generated controller and test code and following the examples we saw there. This situation can't last forever, though, and we'll open this chapter with a couple of additions to the site that bring us face-to-face with our Ruby limitations.

4.1.1 A title helper

When we last saw our new application, we had just updated our mostly static pages to use Rails layouts to eliminate duplication in our views ([Listing 4.1](#)).

Listing 4.1. The sample application site layout.

app/views/layouts/application.html.erb



```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
  </head>
  <body>
    <%= yield %>
  </body>
</html>
```

This layout works well, but there's one part that could use a little polish. Recall that the title line
Ruby on Rails Tutorial Sample App | <%= @title %>

relies on the definition of `@title` in the actions, such as

```
class PagesController < ApplicationController

  def home
    @title = "Home"
  end

  .
  .
  .
```

But what if we don't define an `@title` variable? It's a good convention to have a *base title* we use on


every page, with an optional variable title if we want to be more specific. We’ve *almost* achieved that with our current layout, with one wrinkle: as you can see if you delete the `@title` assignment in one of the actions, in the absence of an `@title` variable the title appears as follows:

Ruby on Rails Tutorial Sample App |

In other words, there’s a suitable base title, but there’s also a trailing vertical bar character `|` at the end of the title.

One common way to handle this case is to define a *helper*, which is a function designed for use in views. Let’s define a `title` helper that returns a base title, “Ruby on Rails Tutorial Sample App”, if no `@title` variable is defined, and adds a vertical bar followed by the variable title if `@title` is defined ([Listing 4.2](#)).¹

Listing 4.2. Defining a `title` helper.

`app/helpers/application_helper.rb` 

```
module ApplicationHelper

  # Return a title on a per-page basis.
  def title
    base_title = "Ruby on Rails Tutorial Sample App"
    if @title.nil?
      base_title
    else
      "#{base_title} | #{@title}"
    end
  end
end
```

This may look fairly simple to the eyes of an experienced Rails developer, but it’s *full* of new Ruby ideas: modules, comments, local variable assignment, booleans, control flow, string interpolation, and return values. We’ll cover each of these ideas in this chapter.

Now that we have a helper, we can use it to simplify our layout by replacing


```
<title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
```

with

```
<title><%= title %></title>
```

as seen in [Listing 4.3](#). Note in particular the switch from the instance variable `@title` to the helper method `title` (without the `@` sign). Using Autotest or `spec spec/`, you should verify that the tests from [Chapter 3](#) still pass.

Listing 4.3. The sample application site layout.

`app/views/layouts/application.html.erb` 

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
```

```

</head>
<body>
  <%= yield %>
</body>
</html>

```

4.1.2 Cascading Style Sheets

There's a second addition to our site that seems simple but adds several new Ruby concepts: including style sheets into our site layout. Though this is a book in web development, not web design, we'll be using cascading style sheets (CSS) to give the sample application some minimal styling, and we'll use the [Blueprint CSS](#) framework as a foundation for that styling.

To get started, [download the latest Blueprint CSS](#). (For simplicity, I'll assume you download Blueprint to a **Downloads** directory, but use whichever directory is most convenient.) Using either the command line or a graphical tool, copy the Blueprint CSS directory **blueprint** into the **public/stylesheets** directory, a special directory where Rails keeps stylesheets. On my Mac, the commands looked like this, but your details may differ:

```

$ cp -r ~/Downloads/joshuaclayton-blueprint-css-016c911/blueprint \
> public/stylesheets/

```

Here **cp** is the Unix copy command, and the **-r** flag copies recursively (needed for copying directories). (As mentioned briefly in [Section 3.2.1.2](#), the tilde **~** means “home directory” in Unix.)

Once you have the stylesheets in the proper directory, Rails provides a helper for including them on our pages using Embedded Ruby ([Listing 4.4](#)).

Listing 4.4. Adding stylesheets to the sample application layout.

`app/views/layouts/application.html.erb`



```

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
    <%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
    <%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
  </head>
  <body>
    <%= yield %>
  </body>
</html>

```

Let's focus on the new lines:

```

<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>

```

These use the built-in Rails helper `stylesheet_link_tag`, which you can read more about at the [Rails API.2](#) The first `stylesheet_link_tag` line includes the stylesheet `blueprint/screen.css` for screens (e.g., computer monitors), and the second includes

`blueprint/print.css` for printing. (The helper automatically appends the `.css` extension to the filenames if absent, so I've left it off for brevity.) As with the `title` helper, to an experienced Rails developer these lines look simple, but there are at least four new Ruby ideas: built-in Rails methods, method invocation with missing parentheses, symbols, and hashes. In this chapter we'll cover these new ideas as well. (We'll see the HTML produced by these stylesheets includes in [Listing 4.6](#) of [Section 4.3.4](#).)

By the way, with the new stylesheets the site doesn't look much different than before, but it's a start ([Figure 4.1](#)). We'll build on this foundation starting in [Chapter 5.3](#)



Figure 4.1: The Home page with the new Blueprint stylesheets. ([full size](#))

[4.2 Strings and methods](#)

Our principal tool for learning Ruby will be the *Rails console*, which is a command-line tool for interacting with Rails applications. The console itself is built on top of interactive Ruby (`irb`), and thus has access to the full power of Ruby. (As we'll see in [Section 4.4.4](#), the console also has access to the Rails environment.) Start the console at the command line as follows:

```
$ script/console
Loading development environment (Rails 2.3.8)
>>
```

By default, the console starts in a *development environment*, which is one of three separate environments defined by Rails (the others are *test* and *production*). This distinction won't be important in this chapter; we'll learn more about environments in [Section 6.3.1](#).

The console is a great learning tool, and you should feel free to explore—don't worry, you (probably) won't break anything. When using the console, type Ctrl-C if you get stuck, or Ctrl-D to exit the console altogether.

Throughout the rest of this chapter, you might find it helpful to consult the [Ruby API.4](#) It's packed (perhaps even *too* packed) with information; for example, to learn more about Ruby strings you can look at the [Ruby API entry for the `String` class](#).

[4.2.1 Comments](#)

Ruby *comments* start with the pound sign `#` and extend to the end of the line. Ruby (and hence Rails) ignores comments, but they are useful for human readers (including, often, the original author!). In the code

```
# Return a title on a per-page basis.
def title
  .
  .
  .
```

the first line is a comment indicating the purpose of the subsequent function definition.

You don't ordinarily include comments in console sessions, but for instructional purposes I'll include some comments in what follows, like this:

```
>> 17 + 42    # Integer addition
=> 59
```

If you follow along in this section typing or copying-and-pasting commands into your own console, you can of course omit the comments if you like; the console will ignore them in any case.

4.2.2 Strings

Strings are probably the most important data structure for web applications, since web pages ultimately consist of strings of characters sent from the server to the browser. Let's get started exploring strings with the console:

```
>> ""          # An empty string
=> ""
>> "foo"       # A nonempty string
=> "foo"
```

These are *string literals* (also, amusingly, called *literal strings*), created using the double quote character ". The console prints the result of evaluating each line, which in the case of a string literal is just the string itself.

We can also concatenate strings with the + operator:

```
>> "foo" + "bar"    # String concatenation
=> "foobar"
```

Here the result of evaluating "foo" plus "bar" is the string "foobar".[5](#)

Another way to build up strings is via *interpolation* using the special syntax `#{}:`[6](#)

```
>> first_name = "Michael"    # Variable assignment
=> "Michael"
>> "#{first_name} Hartl"    # String interpolation
=> "Michael Hartl"
```

Here we've *assigned* the value "Michael" to the variable `first_name` and then interpolated it into the string `"#{first_name} Hartl"`. We could also assign both strings a variable name:

```
>> first_name = "Michael"
=> "Michael"
>> last_name = "Hartl"
=> "Hartl"
>> first_name + " " + last_name    # Concatenation, with a space in between
=> "Michael Hartl"
>> "#{first_name} #{last_name}"    # The equivalent interpolation
=> "Michael Hartl"
```

Note that the final two expressions are equivalent, but I prefer the interpolated version; having to add the single space " " seems a bit awkward.

Printing

To *print* a string, the most commonly used Ruby function is `puts` (pronounced “put ess”, for “put string”):

```
>> puts "foo"      # put string
foo
=> nil
```

The `puts` method operates as a *side-effect*: the expression `puts "foo"` prints the string to the screen and then returns literally nothing: `nil` is a special Ruby value for “nothing at all”. (In what follows, I’ll sometimes suppress the `=> nil` part for simplicity.)

Using `puts` automatically appends a newline character `\n` to the output; the related `print` method does not:

```
>> print "foo"      # print string (same as puts, but without the newline)
foo=> nil
>> print "foo\n"    # Same as puts "foo"
foo
=> nil
```

Single-quoted strings

All the examples so far have used *double-quoted strings*, but Ruby also supports *single-quoted strings*. For many uses, the two types of strings are effectively identical:

```
>> 'foo'            # A single-quoted string
=> "foo"
>> 'foo' + 'bar'
=> "foobar"
```

There’s an important difference, though; Ruby won’t interpolate into single-quoted strings:

```
>> '#{foo} bar'     # Single-quoted strings don't allow interpolation
=> "\#{foo} bar"
```

Note how the console returns values using double-quoted strings, which requires a backslash to *escape* characters like `#`.

If double-quoted strings can do everything that single-quoted strings can do, and interpolate to boot, what’s the point of single-quoted strings? They are often useful because they are truly literal, and contain exactly the characters you type. For example, the “backslash” character is special on most systems, as in the literal newline `\n`. If you want a variable to contain a literal backslash, single quotes make it easier:

```
>> '\n'             # A literal 'backslash n' combination
=> "\\n"
```

As with the `#` character in our previous example, Ruby needs to escape the backslash with an additional backslash; inside double-quoted strings, a literal backslash is represented with *two* backslashes.

For a small example like this, there’s not much savings, but if there are lots of things to escape it can be a real help:


```
>> 'Newlines (\n) and tabs (\t) both use the backslash character \.'
=> "Newlines (\\n) and tabs (\\t) both use the backslash character \\. "
```

4.2.3 Objects and message passing

Everything in Ruby, including strings and even `nil`, is an *object*. We'll see the technical meaning of this in [Section 4.4.2](#), but I don't think anyone ever understood objects by reading the definition in a book; you have to build up your intuition for objects by seeing lots of examples.

It's easier to describe what objects *do*, which is respond to messages. An object like a string, for example, can respond to the message `length`, which returns the number of characters in the string:

```
>> "foobar".length          # Passing the "length" message to a string
=> 6
```

Typically, the messages that get passed to objects are *methods*, which are functions defined on those objects.⁷ Strings also respond to the `empty?` method:

```
>> "foobar".empty?
=> false
>> "".empty?
=> true
```

Note the question mark at the end of the `empty?` method. This is a Ruby convention indicating that the return value is *boolean*: `true` or `false`. Booleans are especially useful for *control flow*:

```
>> s = "foobar"
>> if s.empty?
>>   "The string is empty"
>> else
>>   "The string is nonempty"
>> end
=> "The string is nonempty"
```

Booleans can also be combined using the `&&` (“and”), `||` (“or”), and `!` (“not”) operators:

```
>> x = "foo"
=> "foo"
>> y = ""
=> ""
>> puts "Both strings are empty" if x.empty? && y.empty?
=> nil
>> puts "One of the strings is empty" if x.empty? || y.empty?
"One of the strings is empty"
=> nil
>> puts "x is not empty" if !x.empty?
=> "x is not empty"
```

Since everything in Ruby is an object, it follows that `nil` is an object, so it too can respond to methods. One example is the `to_s` method that can convert virtually any object to a string:

```
>> nil.to_s
""
```

This certainly appears to be an empty string, as we can verify by *chaining* the messages we pass to `nil`:

```
>> nil.empty?
NoMethodError: You have a nil object when you didn't expect it!
You might have expected an instance of Array.
The error occurred while evaluating nil.empty?
>> nil.to_s.empty?      # Message chaining
true
```

We see here that the `nil` object doesn't itself respond to the `empty?` method, but `nil.to_s` does.

There's a special method for testing for `nil`-ness, which you might be able to guess:

```
>> "foo".nil?
=> false
>> "".nil?
=> false
>> nil.nil?
=> true
```

If you look back at [Listing 4.2](#), you'll see that the `title` helper tests to see if `@title` is `nil` using the `nil?` method. This is a hint that there's something special about instance variables (variables with an `@` sign), which can best be understood by contrasting them with ordinary variables. For example, suppose we enter `title` and `@title` variables at the console without defining them first:

```
>> title      # Oops! We haven't defined a title variable.
NameError: undefined local variable or method `title'
>> @title     # An instance variable in the console
=> nil
>> puts "There is no such instance variable." if @title.nil?
There is no such instance variable.
=> nil
>> "#{@title}" # Interpolating @title when it's nil
=> ""
```

You can see from this example that Ruby complains if we try to evaluate an undefined local variable, but issues no such complaint for an instance variable; instead, instance variables are `nil` if not defined. This also explains why the code

```
Ruby on Rails Tutorial Sample App | <%= @title %>
```

becomes

```
Ruby on Rails Tutorial Sample App |
```

when `@title` is `nil`: Embedded Ruby inserts the string corresponding to the given variable, and the string corresponding to `nil` is `""`.

The last example also shows an alternate use of the `if` keyword: Ruby allows you to write a statement that is evaluated only if the statement following `if` is true. There's a complementary `unless` keyword that works the same way:

```
>> string = "foobar"
>> puts "The string '#{string}' is nonempty." unless string.empty?
The string 'foobar' is nonempty.
=> nil
```

It's worth noting that the `nil` object is special, in that it is the *only* Ruby object that is false in a boolean context, apart from `false` itself:

```
>> if nil
>>   true
>> else
>>   false      # nil is false
>> end
=> false
```

In particular, all other Ruby objects are *true*, even 0:

```
>> if 0
>>   true      # 0 (and everything other than nil and false itself) is true
>> else
>>   false
>> end
=> true
```

4.2.4 Method definitions

The console allows us to define methods the same way we did with the `home` action from [Listing 3.3](#) or the `title` helper from [Listing 4.2](#). (Defining methods in the console is a bit cumbersome, and ordinarily you would use a file, but it's convenient for demonstration purposes.) For example, let's define a function `string_message` that takes a single *argument* and returns a message based on whether the argument is empty or not:

```
>> def string_message(string)
>>   if string.empty?
>>     "It's an empty string!"
>>   else
>>     "The string is nonempty."
>>   end
>> end
=> nil
>> puts string_message("")
It's an empty string!
>> puts string_message("foobar")
The string is nonempty.
```

Note that Ruby functions have an *implicit return*, meaning they return the last statement evaluated—in this case, one of the two message strings, depending on whether the method's argument `string` is empty or not. Ruby also has an explicit return option; the following function is equivalent to the one above:

```
>> def string_message(string)
>>   return "It's an empty string!" if string.empty?
>>   return "The string is nonempty."
>> end
```

The alert reader might notice at this point that the second `return` here is actually unnecessary—being the last expression in the function, the string `"The string is nonempty."` will be returned regardless of the `return` keyword, but using `return` in both places has a pleasing symmetry to it.

4.2.5 Back to the `title` helper

We are now in a position to understand the `title` helper from [Listing 4.2:8](#)

```
module ApplicationHelper
  # Return a title on a per-page basis.          # Documentation comment
  def title                                       # Method definition
    base_title = "Ruby on Rails Tutorial Sample App" # Variable assignment
    if @title.nil?                               # Boolean test for nil
      base_title                                 # Implicit return
    else
      "#{base_title} | #{@title}"                # String interpolation
    end
  end
end
end
```

These elements—function definition, variable assignment, boolean tests, control flow, and string interpolation—come together to make a compact helper method for use in our site layout. The final element is `module ApplicationHelper`: code in Ruby modules can be *mixed in* to Ruby classes. When writing ordinary Ruby, you often write modules and include them explicitly yourself, but in this case Rails handles the inclusion automatically for us. The result is that the `title` method is [automagically](#) available in all our views.

4.3 Other data structures

Though web apps are ultimately about strings, actually *making* those strings requires using other data structures as well. In this section we'll learn about some Ruby data structures important for writing Rails applications.

4.3.1 Arrays and ranges

An array is just a list of elements in a particular order. We haven't discussed arrays yet in *Rails Tutorial*, but understanding them gives a good foundation for understanding hashes ([Section 4.3.3](#)) and for aspects of Rails data modeling (such as the `has_many` association seen in [Section 2.3.3](#) and covered more in [Section 11.1.2](#)).

So far we've spent a lot of time understanding strings, and there's a natural way to get from strings to arrays using the `split` method:

```
>> "foo bar    baz".split      # Split a string into a three-element array
=> ["foo", "bar", "baz"]
```

The result of this operation is an array of three strings. By default, `split` divides a string into an array by splitting on whitespace, but you can split on nearly anything else:

```
>> "fooxbarxbazx".split('x')
=> ["foo", "bar", "baz"]
```

As is conventional in most computer languages, Ruby arrays are *zero-offset*, which means that the first element in the array has index 0, the second has index 1, and so on:

```
>> a = [42, 8, 17]
=> [42, 8, 17]
>> a[0]           # Ruby uses square brackets for array access.
=> 42
>> a[1]
=> 8
>> a[2]
=> 17
>> a[-1]          # Indices can even be negative!
=> 17
```

We see here that Ruby uses square brackets to access array elements. In addition to this bracket notation, Ruby offers synonyms for some commonly accessed elements:[9](#)

```
>> a               # Just a reminder of what 'a' is
=> [42, 8, 17]
>> a.first
=> 42
>> a.second
=> 8
>> a.last
=> 17
>> a.last == a[-1] # Comparison using ==
=> true
```

This last line introduces the equality comparison operator `==`, which Ruby shares with many other languages, along with the associated `!=` (“not equal”), etc.:

```
>> x = a.length    # Like strings, arrays respond to the 'length' method.
=> 3
>> x == 3
=> true
>> x == 1
=> false
>> x != 1
=> true
>> x >= 1
=> true
>> x < 1
=> false
```

In addition to `length` (seen in the first line above), arrays respond to a wealth of other methods:

```
>> a.sort
=> [8, 17, 42]
>> a.reverse
=> [17, 8, 42]
>> a.shuffle
=> [17, 42, 8]
```

By the way, the `shuffle` method is available only on Ruby 1.8.7 or later; if you’re using 1.8.6, you can use

```
>> a.sort_by { rand }
```

instead.

You can also add to arrays with the “push” operator, `<<`:

```
>> a << 7                # Pushing 7 onto an array
[42, 8, 17, 7]
>> a << "foo" << "bar"    # Chaining array pushes
[42, 8, 17, 7, "foo", "bar"]
```

This last example shows that you can chain pushes together, and also that, unlike arrays in many other languages, Ruby arrays can contain a mixture of different types (in this case, integers and strings).

Before we saw `split` convert a string to an array. We can also go the other way with the `join` method:

```
>> a
[42, 8, 17, 7, "foo", "bar"]
>> a.join                # Join on nothing
=> "428177foobar"
>> a.join(',')           # Join on comma-space
=> "42, 8, 17, 7, foo, bar"
```

Closely related to arrays are *ranges*, which can probably most easily be understood by converting them to arrays using the `to_a` method:

```
>> 0..9
=> 0..9
>> 0..9.to_a             # Oops, call to_a on 9
ArgumentError: bad value for range
>> (0..9).to_a           # Use parentheses to call to_a on the range
=> [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Though `0..9` is a valid range, the second expression above shows that we need to add parentheses to call a method on it.

Ranges are useful for pulling out array elements:

```
>> a = %w[foo bar baz quux]    # Use %w to make a string array
["foo", "bar", "baz", "quux"]
>> a[0..2]
=> ["foo", "bar", "baz"]
```

Ranges also work with characters:

```
>> ('a'..'e').to_a
=> ["a", "b", "c", "d", "e"]
```

4.3.2 Blocks

Both arrays and ranges respond to a host of methods that accept *blocks*, which are simultaneously one of Ruby’s most powerful and most confusing features:

```
>> (1..5).each { |i| puts 2 * i }
2
4
6
```

8
10
=> 1..5

This code calls the `each` method on the range `(1..5)` and passes it the block `{ |i| puts 2 * i }`. The vertical bars around the variable name in `|i|` are Ruby syntax for a block variable, and it's up to the method to know what to do with the block; in this case, the range's `each` method can handle a block with a single local variable, which we've called `i`, and it just executes the block for each value in the range.

Curly braces are one way to indicate a block, but there is a second way as well:

```
>> (1..5).each do |i|
?>   puts 2 * i
>> end
2
4
6
8
10
=> 1..5
```

Blocks can be more than one line, and often are. In *Rails Tutorial* we'll follow the common convention of using curly braces only for short one-line blocks and the `do . . end` syntax for longer one-liners and for multi-line blocks:

```
>> (1..5).each do |number|
?>   puts 2 * number
>>   puts '--'
>> end
2
--
4
--
6
--
8
--
10
--
=> 1..5
```

Here I've used `number` in place of `i` just to emphasize that any variable name will do.

Unless you already have a substantial programming background, there is no shortcut to understanding blocks; you just have to see them a lot, and eventually you'll get used to them.¹⁰ Luckily, humans are quite good at making generalizations from concrete examples; here are a few more, including a couple using the `map` method:

```
>> 3.times { puts "Betelgeuse!" } # 3.times takes a block with no variables.
"Betelgeuse!"
"Betelgeuse!"
"Betelgeuse!"
=> 3
>> (1..5).map { |i| i**2 } # The ** notation is for 'power'.
=> [1, 4, 9, 16, 25]
>> %w[a b c] # Recall that %w makes string arrays.
```

```
=> ["a", "b", "c"]
>> %w[a b c].map { |char| char.upcase }
=> ["A", "B", "C"]
```

As you can see, the `map` method returns the result of applying the given block to each element in the array or range.

By the way, we're now in a position to understand the line of Ruby I threw into [Section 1.4.4](#) to generate random subdomains:[11](#)

```
('a'..'z').to_a.shuffle[0..7].join
```

Let's build it up step-by-step:

```
>> ('a'..'z').to_a                # An alphabet array
=> ["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l", "m", "n", "o",
    "p", "q", "r", "s", "t", "u", "v", "w", "x", "y", "z"]
>> ('a'..'z').to_a.shuffle        # Shuffle it!
=> ["c", "g", "l", "k", "h", "z", "s", "i", "n", "d", "y", "u", "t", "j", "q",
    "b", "r", "o", "f", "e", "w", "v", "m", "a", "x", "p"]
>> ('a'..'z').to_a.shuffle[0..7]  # Pull out the first eight elements.
=> ["f", "w", "i", "a", "h", "p", "c", "x"]
>> ('a'..'z').to_a.shuffle[0..7].join # Join them together to make one string.
=> "mznpbyuj"
```

4.3.3 Hashes and symbols

Hashes are essentially a generalization of arrays: you can think of hashes as basically like arrays, but not limited to integer indices. (In fact, some languages, especially Perl, call hashes *associative arrays* for this reason.) Instead, hash indices, or *keys*, can be almost any object. For example, we can use strings as keys:

```
>> user = {}                    # {} is an empty hash
=> {}
>> user["first_name"] = "Michael" # Key "first_name", value "Michael"
=> "Michael"
>> user["last_name"] = "Hartl"   # Key "last_name", value "Hartl"
=> "Hartl"
>> user["first_name"]           # Element access is like arrays
=> "Michael"
>> user                         # A literal representation of the hash
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}
```

Hashes are indicated with curly braces containing key-value pairs; a pair of braces with no key-value pairs—i.e., `{}`—is an empty hash. It's important to note that the curly braces for hashes have nothing to do with the curly braces for blocks. (Yes, this can be confusing.) Though hashes resemble arrays, one important difference is that hashes don't generally guarantee keeping their elements in a particular order.[12](#) If order matters, use an array.

Instead of defining hashes one item at a time using square brackets, it's easy to use their literal representation:

```
>> user = { "first_name" => "Michael", "last_name" => "Hartl" }
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}
```


Here I've used the usual Ruby convention of putting an extra space at the two ends of the hash—a convention ignored by the console output. (Don't ask me why the spaces are conventional; probably some early influential Ruby programmer liked the look of the extra spaces, and the convention stuck.)

So far we've used strings as hash keys, but in Rails it is much more common to use *symbols* instead. Symbols look kind of like strings, but prefixed with a colon instead of surrounded by quotes. For example, `:name` is a symbol. You can think of symbols as basically strings without all the extra baggage:[13](#)

```
>> "name".length
=> 4
>> :name.length
NoMethodError: undefined method `length' for :name:Symbol
>> "foobar".reverse
=> "raboof"
>> :foobar.reverse
NoMethodError: undefined method `reverse' for :foobar:Symbol
```

Symbols are a special Ruby data type shared with very few other languages, so they may seem weird at first, but Rails uses them a lot, so you'll get used to them fast.

In terms of symbols as hash keys, we can define a `user` hash as follows:

```
>> user = { :name => "Michael Hartl", :email => "michael@example.com" }
=> {:name=>"Michael Hartl", :email=>"michael@example.com"}
>> user[:name]           # Access the value corresponding to :name.
=> "Michael Hartl"
>> user[:password]       # Access the value of an undefined key.
=> nil
```

We see here from the last example that the hash value for an undefined key is simply `nil`.

Hash values can be virtually anything, even other hashes, as seen in [Listing 4.5](#).

Listing 4.5. Nested hashes.

```
>> params = {}           # Define a hash called 'params' (short for 'parameters').
=> {}
>> params[:user] = { :name => "Michael Hartl", :email => "mhartl@example.com" }
=> {:name=>"Michael Hartl", :email=>"mhartl@example.com"}
>> params
=> {:user=>{:name=>"Michael Hartl", :email=>"mhartl@example.com"}}
>> params[:user][:email]
=> "mhartl@example.com"
```

These sorts of hashes-of-hashes, or *nested hashes*, are heavily used by Rails, as we'll see starting in [Section 8.2](#).

As with arrays and ranges, hashes respond to the `each` method. For example, consider a hash named `flash` with keys for two conditions, `:success` and `:error`:

```
>> flash = { :success => "It worked!", :error => "It failed. :-( " }
=> {:success=>"It worked!", :error=>"It failed. :-( "}
>> flash.each do |key, value|
?>   puts "Key #{key.inspect} has value #{value.inspect}"
>> end
Key :success has value "It worked!"
```

Key :error has value "It failed. :-("

Note that, while the `each` method for arrays takes a block with only one variable, `each` for hashes takes two, a *key* and a *value*. Thus, the `each` method for a hash iterates through the hash one key-value *pair* at a time.

The last example uses the useful `inspect` method, which returns a string with a literal representation of the object it's called on:

```
>> puts flash                # Put the flash hash as a string (with ugly results).
successIt worked!errorIt failed. :-(
>> puts flash.inspect        # Put the flash hash as a pretty string
{:success=>"It worked!", :error=>"It failed. :-( "}
>> puts :name, :name.inspect
name
:name
>> puts "It worked!", "It worked!".inspect
It worked!
"It worked!"
```

By the way, using `inspect` to print an object is common enough that there's a shortcut for it, the `p` function:

```
>> p flash                    # Same as 'puts flash.inspect'
{:success=>"It worked!", :error=>"It failed. :-( "} 
```

4.3.4 CSS revisited

It's time now to revisit the lines from [Listing 4.4](#) used in the layout to include the cascading style sheets:

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
```

We are now nearly in a position to understand this. As mentioned briefly in [Section 4.1.2](#), Rails defines a special function to include stylesheets, and

```
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

is a call to this function. But there are two mysteries. First, where are the parentheses? In Ruby, they are optional; these two lines are equivalent:

```
# Parentheses on function calls are optional.
stylesheet_link_tag('blueprint/screen', :media => 'screen')
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

Second, the `:media` argument sure looks like a hash, but where are the curly braces? When hashes are the *last* argument in a function call, the curly braces are optional; these two lines are equivalent:

```
# Curly braces on final hash arguments are optional.
stylesheet_link_tag 'blueprint/screen', { :media => 'screen' }
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

So, we see now that each of the lines

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
```

calls the `stylesheet_link_tag` function with two arguments: a string, indicating the path to the stylesheet, and a hash, indicating the media type (`'screen'` for the computer screen and `'print'` for a printed version). Because of the `<%= %>` brackets, the results are inserted into the template by ERb, and if you view the source of the page in your browser you should see the HTML needed to include a stylesheet ([Listing 4.6](#)).¹⁴

Listing 4.6. The HTML source produced by the CSS includes.



```
<link href="/stylesheets/blueprint/screen.css" media="screen" rel="stylesheet"
type="text/css" />
<link href="/stylesheets/blueprint/print.css" media="print" rel="stylesheet"
type="text/css" />
```

4.4 Ruby classes

We've said before that everything in Ruby is an object, and in this section we'll finally get to define some of our own. Ruby, like many object-oriented languages, uses *classes* to organize methods; these classes are then *instantiated* to create objects. If you're new to object-oriented programming, this may sound like gibberish, so let's look at some concrete examples.

4.4.1 Constructors

We've seen lots of examples of using classes to instantiate objects, but we have yet to do so explicitly. For example, we instantiated a string using the double quote characters, which is a *literal constructor* for strings:

```
>> s = "foobar"          # A literal constructor for strings using double quotes
=> "foobar"
>> s.class
=> String
```

We see here that strings respond to the method `class`, and simply return the class they belong to.

Instead of using a literal constructor, we can use the equivalent *named constructor*, which involves calling the `new` method on the class name:

```
>> s = String.new("foobar") # A named constructor for a string
=> "foobar"
>> s.class
=> String
>> s == "foobar"
=> true
```

This is equivalent to the literal constructor, but it's more explicit about what we're doing.


```
"#{@name} <#{@email}>"
end
```

Let's fire up the console, **require** the example user code, and take our User class out for a spin:

```
>> require 'example_user'      # This is how you load the example_user code.
=> ["User"]
>> example = User.new
=> #<User:0x224ceec @email=nil, @name=nil>
>> example.name                 # nil since attributes[:name] is nil
=> nil
>> example.name = "Example User"    # Assign a non-nil name
=> "Example User"
>> example.email = "user@example.com" # and a non-nil email address
=> "user@example.com"
>> example.formatted_email
=> "Example User <user@example.com>"
```

This code creates an empty example user and then fills in the name and email address by assigning directly to the corresponding attributes (assignments made possible by the `attr_accessor` line in [Listing 4.8](#)). When we write

```
example.name = "Example User"
```

Ruby is setting the `@name` variable to "Example User" (and similarly for the `email` attribute), which we then use in the `formatted_email` method.

Recalling from [Section 4.3.4](#) we can omit the curly braces for final hash arguments, we can create another user by passing a hash to the `initialize` method to create a user with pre-defined attributes:

```
>> user = User.new(:name => "Michael Hartl", :email => "mhartl@example.com")
=> #<User:0x225167c @email="mhartl@example.com", @name="Michael Hartl">
>> user.formatted_email
=> "Michael Hartl <mhartl@example.com>"
```

We will see starting in [Chapter 8](#) that initializing objects using a hash argument is common in Rails applications.

4.5 Exercises

1. Using [Listing 4.9](#) as a guide, combine the `split`, `shuffle`, and `join` methods to write a function that shuffles the letters in a given string.
2. Using [Listing 4.10](#) as a guide, add a `shuffle` method to the `String` class.
3. Create three hashes called `person1`, `person2`, and `person3`, with first and last names under the keys `:first` and `:last`. Then create a `params` hash so that `params[:father]` is `person1`, `params[:mother]` is `person2`, and `params[:child]` is `person3`. Verify that, for example, `params[:father][:first]` has the right value.
4. Learn about the `Hash` method `merge`.

Listing 4.9. Skeleton for a string shuffle function.

```
>> def string_shuffle(s)
>>   s.split('').??.?
>> end
=> nil
>> string_shuffle("foobar")
```

Listing 4.10. Skeleton for a `shuffle` method attached to the `String` class.

```
>> class String
>>   def shuffle
>>     self.split('').??.?
>>   end
>> end
=> nil
>> "foobar".shuffle
```

[« Chapter 3 Mostly static pages Chapter 5 Filling in the layout »](#)

1. If a helper is specific to a particular controller, you should put it in the corresponding helper file; for example, helpers for the Pages controller generally go in `app/helpers/pages_helper.rb`. In our case, we expect the `title` helper to be used on all the site's pages, and Rails has a special helper file for this case: `app/helpers/application_helper.rb`. [↑](#)
2. I've linked my favorite Rails API source, [RailsBrain](#). Its API version is currently a little out of date, but it doesn't usually matter much. If you think it might matter, you can use the [official Rails API](#), but I prefer the RailsBrain interface. Incidentally, "API" stands for [application programming interface](#). [↑](#)
3. If you're impatient, feel free to check out the [Blueprint CSS Quickstart tutorial](#). [↑](#)
4. I've linked to [RubyBrain](#), sister site to [RailsBrain](#). [↑](#)
5. For more on the origins of "foo" and "bar"—and, in particular, the possible *non*-relation of "foobar" to "FUBAR"—see the [Jargon File entry on "foo"](#). [↑](#)
6. Programmers familiar with Perl or PHP should compare this to the automatic interpolation of dollar sign variables in expressions like "`foo $bar`". [↑](#)
7. Apologies in advance for switching haphazardly between *function* and *method* throughout this chapter; in Ruby, they're the same thing: all methods are functions, and all functions are methods, because everything is an object. [↑](#)
8. Well, there will still be *one* thing left that we don't understand, which is how Rails ties this all together: mapping URLs to actions, making the `title` helper available in views, etc. This is an interesting subject, and I encourage you to investigate it further, but knowing exactly *how* Rails works is not necessary to *using* Rails. (For a deeper understanding, I recommend [The Rails Way](#) by Obie Fernandez.) [↑](#)
9. The `second` method used here isn't currently part of Ruby itself, but rather is added by Rails. It works in this case because the Rails console automatically includes the Rails extensions to Ruby. [↑](#)
10. Programming experts, on the other hand, might benefit from knowing that blocks are *closures*, which are one-shot anonymous functions with data attached. [↑](#)
11. Recall that `shuffle` isn't available in Ruby 1.8.6; you have to use `sort_by { rand }` instead. [↑](#)
12. Apparently Ruby 1.9 guarantees that hashes keep their elements in the same order entered, but

it would be unwise ever to count on a particular ordering. [↑](#)

13.As a result of having less baggage, symbols are easier to compare to each other; strings need to be compared character by character, while symbols can be compared all in one go. This makes them ideal for use as hash keys. [↑](#)

14.You may see some funky numbers, like ?1257465942, after the CSS filenames. These are inserted by Rails to ensure that browsers reload the CSS when it changes on the server. [↑](#)

15.For more on Ruby classes and the `self` keyword, see the [RailsTips](#) post on [Class and Instance Variables in Ruby](#). [↑](#)

16.For those familiar with JavaScript, this functionality is comparable to using a built-in class prototype object to augment the class. (Thanks to reader [Erik Eldridge](#) for pointing this out.) [↑](#)

17.You don't have to know what each class in this hierarchy does. *I* don't know what they all do, and I've been programming in Ruby on Rails since 2005. This means either that (a) I'm grossly incompetent or (b) you can be a skilled Rails developer without knowing all its innards. I hope for both our sakes that it's the latter. [↑](#)